

Introduction to Database Systems

IDBS – Fall 2024

- Week 2:
- SQL Select
- SQL Joins
- SQL Aggregations

Eleni Tziritza Zacharatou

Readings:
PDBM 7.3

- **Data Definition Language (DDL)**

- Used by the database administrator (DBA) to define the database's data model
- Three common commands:
 - CREATE TABLE, ALTER TABLE, and DROP TABLE

- **Data Manipulation Language (DML)**

- Used by applications and users to retrieve, insert, modify, and delete records
- Four statements:
 - SELECT, INSERT, UPDATE, and DELETE

Today's focus

Database Schema

Running Example

- Coffees(name, manufacturer)
- Coffeehouses(name, address, license)
- Drinkers(name, address, phone)
- Likes(drinker, coffee)
- Sells(coffeehouse, coffee, price)
- Frequents(drinker, coffeehouse)

Note

Keys are underlined



User



declaratively



Query

DBMS Interface

SQL DML -- Select

Readings:
PDBM 1

Rationale & Example

- **SELECT** -- desired attributes
FROM -- one or more relations
WHERE -- condition about records of the relations

- Which coffees are made by Ottolina?

Coffees

name	manufacturer

sql

```
SELECT name FROM Coffees  
WHERE manufacturer = 'Ottolina';
```

- **Result**
 - a relation
 - a single attribute



name
Fortissima
Maracaibo
Buongiorno

The Same but Graphically

name	manufacturer
...	...
...	...
Maracaibo	Ottolina
...	...

```
sql
SELECT name FROM Coffees
WHERE manufacturer = 'Ottolina';
```

filter
Check if Ottolina
If so, project *name*

iterate
add

Note

SQL uses bag semantics

Allows for duplicates

name
Fortissima
Maracaibo

Star in SELECT

- What if the **SELECT** clause has * instead?

Coffees

name	manufacturer

sql

```
SELECT * FROM Coffees  
WHERE manufacturer = 'Ottolina';
```

- This selects all attributes from the **Coffees** relation

name	manufacturer
Fortissima	Ottolina
Maracaibo	Ottolina
Buongiorno	Ottolina

Attribute Renaming

- What If you want the result to have different attribute name?
 - Use then the keyword AS

Coffees	
name	manufacturer

sql

```
SELECT name AS coffee, manufacturer FROM Coffees  
WHERE manufacturer = 'Ottolina';
```

coffee	manufacturer
Fortissima	Ottolina
Maracaibo	Ottolina
Buongiorno	Ottolina

Expressions in SELECT

- Any expression that makes sense can appear as an element of a SELECT clause.

```
SELECT coffeehouse, coffee, price * 0.965 AS priceInYen FROM Sells;
```

Sells

coffeehouse	coffee	price



coffeehouse	coffee	priceInYen
Sue's	Fortissima	342
Joe's	Maracaibo	285
Bob's	Buongiorno	149

Constants as Expressions in SELECT

- Any constant as expression (example with the Likes relation)

Likes

drinker	coffee

sql

```
SELECT drinker,  
       'likes Fortissima' AS whoLikesFort  
FROM Likes  
WHERE coffee = 'Fortissima';
```



drinker	whoLikesFort
Sally	Likes Fortissima
Fred	Likes Fortissima

Live Exercise

- What should the following queries return?
 - Assume the relation Coffees has two records only

coffee	manufacturer
Fortissima	Ottolina
Maracaibo	Ottolina

coffee	manufacturer
Fortissima	Ottolina
Maracaibo	Ottolina

sql

```
SELECT *  
FROM Coffees;
```

Query

42

42

sql

```
SELECT 42 AS Query  
FROM Coffees
```

Query

42

sql

```
SELECT 42 AS Query;
```

Practice

- **Analog has a 50% discount on a second coffee of the same type. For each coffee sold there, show the price of two coffees ('priceOfTwo').**
- **A coffeehouse owner by the name of "Eleni" has passed out with caffeine shock. Write a query to find her home phone number.**

Conditions in WHERE

- **Comparisons** =, <>, <, >, <=, >=.
 - Many other operators that produce boolean-valued results.
- **Boolean operators** AND, OR, NOT.
- **Find the price Joe's coffeehouse charges for Maracaibo in Sells**

Sells

coffeehouse	coffee	price

sql

```
SELECT price
FROM Sells
WHERE coffeehouse = 'Joe's' AND coffee = 'Maracaibo';
```

Patterns

- A condition can compare a string to a pattern by:
 - <Attribute> LIKE <pattern>
 - <Attribute> NOT LIKE <pattern>
- **Pattern** is a quoted string with
 - % = “any string”
 - _ = “any character”
- **Escaping characters**
 - “\” default escape character
 - ESCAPE modifier
- **PostgreSQL supports regular expressions**
 - These are much more expressive!

sql

```
SELECT * FROM Coffeehouses WHERE address LIKE '%1';
```

sql

```
SELECT * FROM Coffeehouses WHERE address LIKE '%!_1' ESCAPE '!';
```

Practice

- A coffeehouse patron by the last name of “Sivertsen” has passed out, again from caffeine shock. Show a query to find his home phone.
- Someone has recommended a coffeehouse that sells a coffee called “Blue”- something, that costs more than 100. Unfortunately, she doesn’t remember the name of the coffeehouse. Find where that coffee is sold at such a high price.

Sorting

- **SQL assumes bags semantics**
 - No order assumption
- **ORDER BY sorts the results**
 - Ascending (ASC) or descending (DESC)

sql

```
SELECT coffeehouse, price  
FROM Sells  
ORDER BY coffeehouse, price DESC;
```

- **When should we use ORDER BY?**



User



declaratively



Query

DBMS Interface

SQL DML -- Joins

Readings:
PDBM 1

Multi-Relation Queries

- Most typically, queries combine data from more than one relation.
- Several relations in the FROM clause.
- Distinguish attributes of the same name by “<relation>.<attribute>”
- Find the coffees liked by at least one person who frequents Joe's.

sql

```
SELECT coffee
FROM Likes
JOIN Frequents
ON Likes.drinker = Frequents.drinker
WHERE Frequents.coffeehouse = 'Joe's';
```

The Same but Graphically

sql

```
SELECT coffee
FROM Likes
JOIN Frequents
ON Likes.drinker = Frequents.drinker
WHERE Frequents.coffeehouse = 'Joe's';
```

Likes

drinker	coffee
...	...
...	...
Sally	Maracaibo
...	...

join

Frequents

drinker	coffeehouse
...	...
...	...
Sally	Joe's
...	...

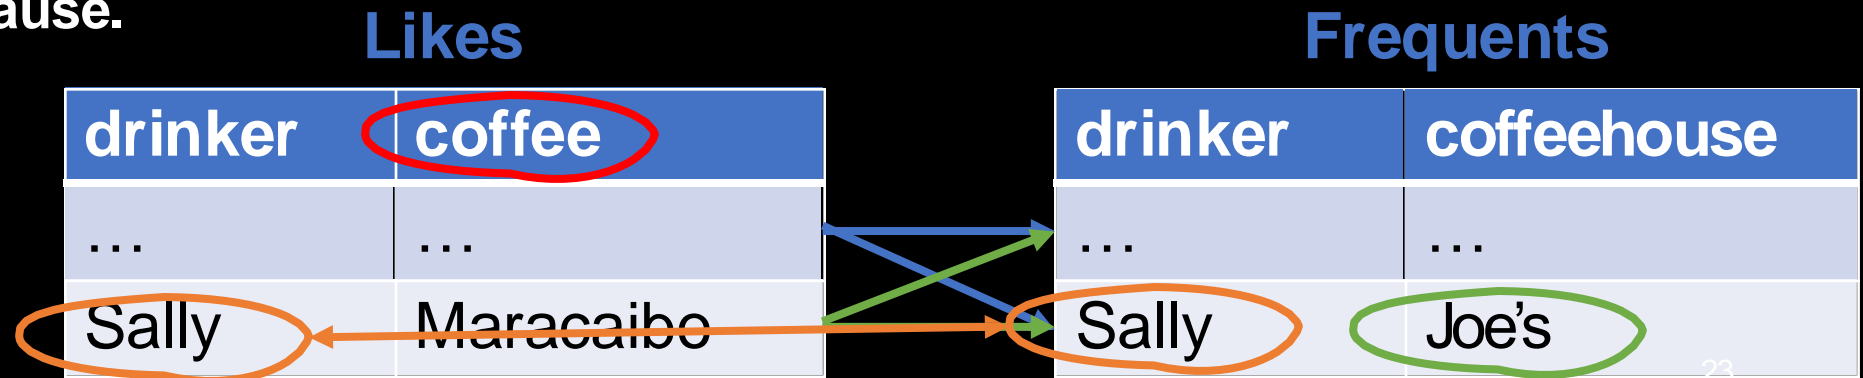
filter

Join Semantics

- **Start with the product of all relations in the FROM clause**
 - Imagine one tuple-variable for each relation in the FROM clause.
 - Think of nested for-loops
- **Apply the selection conditions from the JOIN clauses**
 - Usually $F.\text{foreign_key} = K.\text{key}$
- **Apply the selection condition from the WHERE clause**
- **Apply the selection Project onto the list of attributes and expressions in the SELECT clause.**

sql

```
SELECT coffee
FROM Likes
JOIN Frequents
ON Likes.drinker = Frequents.drinker
WHERE Frequents.coffeehouse = 'Joe's';
```



Renaming Relations

sql

```
SELECT coffee
FROM Likes
JOIN Frequents
ON Likes.drinker = Frequents.drinker
WHERE Frequents.coffeehouse = 'Joe's';
```



sql

```
SELECT L.coffee
FROM Likes L
JOIN Frequents F
ON L.drinker = F.drinker
WHERE F.coffeehouse = 'Joe's';
```

Practice

- For each person that “frequents” some coffeehouse, show the name of the person and the address of the coffeehouse.
- For each person that “frequents” some coffeehouse, show the address of the person and the address of the coffeehouse.
- For each person that “frequents” some coffeehouse, show the name of the person and the address of the coffeehouse.
 - Why don't we need the Drinkers relation?
 - What if we want only the names of drinkers and coffeehouses?
 - What if we used IDs in our design, and wanted the names?
 - What if we want all drinkers, including those that don't frequent any coffeehouse?
 - What if we only want the name of drinkers?

Select + Join

- **SELECT**
FROM
WHERE
JOIN
 - desired attributes
 - one or more relations
 - condition about records of the relations
 - connect records between relations

Duplicate Elimination

- Force the result to be a set with **SELECT DISTINCT ...**
 - Can also do this with GROUP BY – but please don't!
- From **Sells(coffeehouse, coffee, price)**
 - find all the different prices charged for coffees:

Note

Sorting = $O(n \log n)$

sql

```
SELECT DISTINCT price  
FROM Sells;
```

Note

Without DISTINCT, each price would be listed as many times as there were coffeehouse/coffee pairs at that price



User



declaratively



Query

DBMS Interface

SQL DML -- Aggregations

Readings:
PDBM 1

Relation-based Aggregations

- **SUM, AVG, COUNT, MIN, and MAX** can be applied to a column in a **SELECT** clause
 - Produces that aggregation on the column
- **COUNT(*)** counts the number of tuples.
- **Find the average price of Maracaibo**

sql

```
SELECT AVG(price)
FROM Sells
WHERE coffee = 'Maracaibo';
```

Duplicate Elimination

- Use **DISTINCT** inside an aggregation.
- Find the number of different prices charged for Maracaibo

sql

```
SELECT COUNT(DISTINCT price)
FROM Sells
WHERE coffee = 'Maracaibo';
```

Practice

- **Show the highest price of any coffee**

Nulls in Aggregations

- NULL never contributes to a sum, average, or count, and can never be the minimum or maximum of a column.
- But if there are only NULL values in a column
 - The result of the aggregation is NULL
- **Exception:** COUNT of an empty set is 0.

Note

The number of coffeehouses
at a known price!

Sells		
coffeehouse	coffee	price

sql

```
SELECT COUNT(coffeehouse)
FROM Sells
WHERE coffee = 'Maracaibo';
```

sql

```
SELECT COUNT(price)
FROM Sells
WHERE coffee = 'Maracaibo';
```

Grouping-based Aggregations

- **SELECT-FROM-WHERE** expression followed by **GROUP BY**
- **The relation that results from the SELECT-FROM-WHERE is:**
 - grouped according to the values of the attributes in the **GROUP BY** clause
 - any aggregation is applied only within each group
- **How does the system perform GROUP BY?**
 - sort the relation OR hash it!
- **Find the average price for each coffee**

coffee	AVG(price)
Maracaibo	2.33
Fortissima	3.51

sql

```
SELECT coffee, AVG(price)
FROM Sells
GROUP BY coffee;
```

Grouping-based Aggregations

- Find for each drinker the average price of Maracaibo at the coffeehouses they frequent

sql

```
SELECT drinker, AVG(price)
FROM Frequents F
JOIN Sells S ON F.coffeehouse = S.coffeehouse
WHERE coffee = 'Maracaibo'
GROUP BY drinker;
```

Frequents

drinker	coffeehouse

Sells

coffeehouse	coffee	price

SELECT Attributes with Aggregations

- **If any aggregation is used, then each element of the SELECT list must be either:**
 - An attribute on the GROUP BY list
 - Aggregation – COUNT, AVG, MAX, ...
- **All GROUP BY attributes must be in SELECT**
 - **Caveat:** SQL standard vs. Implementation
 - Some systems are more flexible!
 - PostgreSQL allows omitting functionally dependent attributes (see week 6!)

Example Query

Sells

coffeehouse	coffee	price

- For each coffeehouse, show the average price of all coffees in that coffeehouse.

sql

```
SELECT coffeehouse, AVG(price) AS average_price  
FROM Sells  
GROUP BY coffeehouse;
```




coffeehouse	average_price

What About this One?

- Show the price of the most expensive coffee (from any coffeehouse) and the name of the coffeehouse that sells it

```
sql

SELECT coffeehouse, MAX(price)
FROM Sells
GROUP BY coffeehouse;
```



```
sql

SELECT coffeehouse, MAX(price)
FROM Sells
```



```
sql

SELECT coffeehouse, price
FROM Sells
WHERE price = (SELECT MAX(price) FROM Sells);
```

Practice

- For each coffeehouse, show the number of coffees sold in that coffeehouse and the average price.
- For each coffeehouse, show the price of the most expensive coffee sold in that coffeehouse.
- Show the name and price of the least expensive coffee (from any coffeehouse).

HAVING Clause

- **HAVING <condition> may follow a GROUP BY clause.**
 - If so, the condition applies to each group, and groups not satisfying the condition are eliminated.
- **Like WHERE but for groups!**
- **Find the average price of those coffees that are served in at least two coffeehouses**

Sells

coffeehouse	coffee	price

sql

```
SELECT coffee, AVG(price)
FROM Sells
GROUP BY coffee
HAVING COUNT(coffeehouse) > 1
```

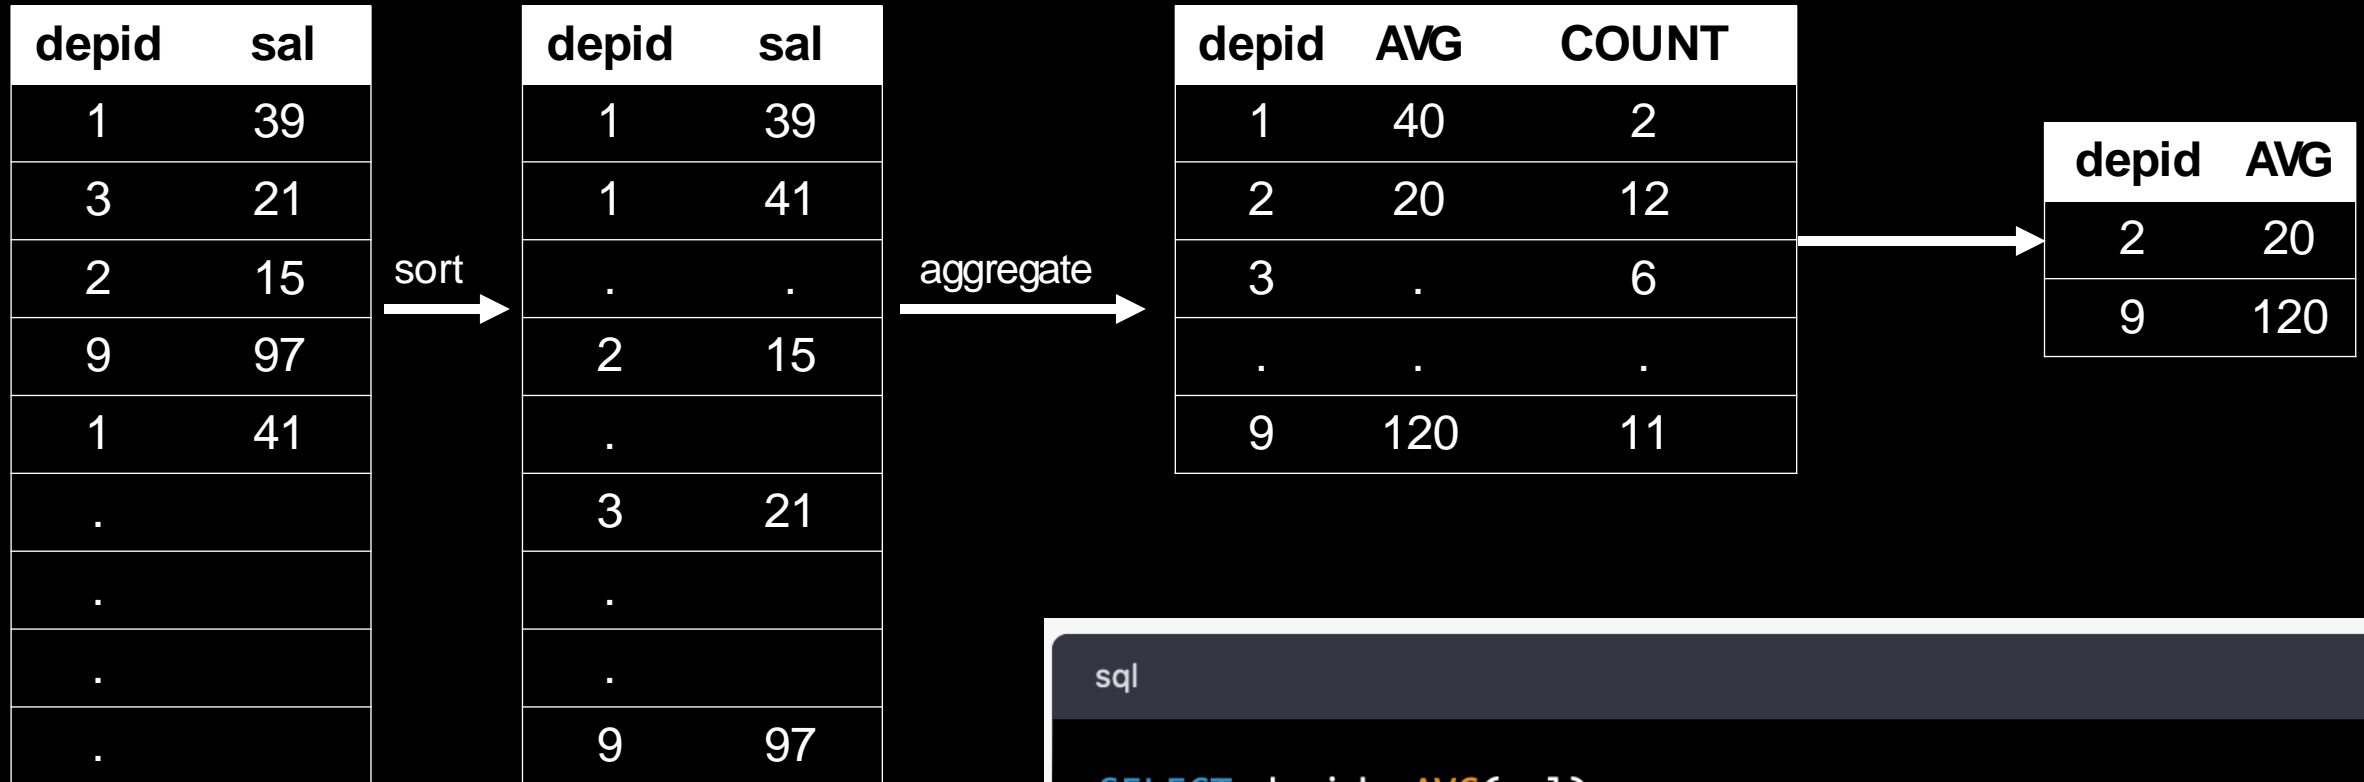
HAVING Conditions

- Anything goes in a sub-query – later...
- Outside sub-queries, they may refer to attributes only if they are either
 - A grouping attribute, or
 - Aggregated(same condition as for SELECT clauses with aggregation)

Practice

- For each coffeehouse that sells more than two coffees, show the number of coffees sold in that coffeehouse and their average price.
- For each coffeehouse, show the number of drinkers that frequent the coffeehouse.
- For each drinker that frequents more than one coffeehouse, show the number of coffeehouses that he/she frequents.


GROUP BY + HAVING Implementation




sql

```
SELECT depid, AVG(sal)
FROM Emp
GROUP BY depid
HAVING COUNT(*) > 10
```

Takeaways

- 
- **SQL universal relational query language**
 - Result set (bag) is a table
 - **1 Block**
 - select, from, where, group by, having, sort by clauses
 - **Joins**
 - Combine data from many tables
 - **Aggregations**
 - Group data
 - **SQL is **code****
 - Treat it accordingly

What is next?

- 
- **Next week: Advanced SQL**
 - **Exercise 2 is out (SQL)**
 - Remember: No submission – but very important to do it!
 - It has some queries that you did not learn how to do yet – just think about potential solutions for now, we will cover them next week.
 - **Homework 1 description is out**
 - SQL queries, some being more complex
 - PDF and SQL script are already available
 - Quiz on LearnIT will open at the end of the week
 - **For help:** Talk to TAs in exercise class
 - **Due Sep 23 at 23:59** – individual submissions!
 - **Remember:** We only accept and give feedback to HWs submitted before the deadline.
 - **Do the exercises before the homework!**