

Introduction to Database Systems

IDBS - Fall 2024

Lecture 4 - SQL Programming and Python

Functions

Triggers & Constraints

SQL and DBMS in Python

Transactions

Readings: PDBM 9.2, 14.1, 14.2.1 and 14.5

Omar Shahbaz Khan

Last Time in IDBS...

-- TODO -> DONE

- ✓ Division
- ✓ JOIN & NULL
 - ✓ Natural Joins, Cross Joins, Self-Joins
 - ✓ NULL
 - ✓ Outer Joins (LEFT, RIGHT, FULL)
- ✓ Set Operations
 - ✓ UNION, INTERSECT, EXCEPT
- ✓ Subqueries (Nested Queries)
 - ✓ =, IN, EXISTS, ALL, ANY
- ✓ Views: Queries as subroutines

Wake Up Task!

Bills DB

People (PID, pName, pGender, pHeight)

Accounts (AID, PID, aDate, aBalance, aOver)

AccountRecords (RID, AID, rDate, rType, rAmount, rBalance)

Bills (BID, PID, bDueDate, bAmount, bIsPaid)

- How many accounts have never been used before?
- How many accounts have a negative account balance and also have bills due?

This Time...

-- TODO

- Database Functions
- Database Triggers & Constraints
- Transactions
- SQL in Python

$$f(x)$$

Functions in SQL

- What is a function?
- Similar to other programming languages methods/functions
- How is it useful in SQL?

Create a new Person

- In the Sports database create a new person
 - Insert a row
- Can be cumbersome to write `INSERT INTO ... VALUES ...` if we want to insert many new rows
- How about a function?

```
INSERT INTO  
People (name, gender, height)  
VALUES ('Terry', 'M', 1.77);
```

Create a new Person

- In the Sports database create a new person
 - Insert a row
- Can be cumbersome to write `INSERT INTO ... VALUES ...` if we want to insert many new rows
- How about a function?

```
DROP FUNCTION IF EXISTS NewPerson;
```

```
CREATE FUNCTION NewPerson (
```

Function
Parameters

```
IN pname VARCHAR(50),  
IN pgender CHAR(1),  
IN pheight FLOAT
```

```
)
```

```
RETURNS INTEGER
```

Return Type

```
AS $$
```

```
BEGIN
```

Insert
Statement

```
INSERT INTO  
Person (name, gender, height)  
VALUES (pname, pgender, pheight);  
RETURN lastval();
```

Return the last value of
the table's SEQUENCE

```
END;
```

```
$$ LANGUAGE plpgsql;
```

Using Functions

- From an SQL Script
- Typically used from an ODBC (Java, Python, C#, ...)

```
SELECT NewPerson('Terry', 'M', 1.77);
```

```
SELECT * FROM NewPerson('Terry', 'M', 1.77);
```

```
DO $$  
    BEGIN  
        PERFORM NewPerson('Terry', 'M', 1.77);  
    END  
$$;
```

Let's make a function (I)

- Assume RecordLog table:

```
RecordLog (  
    peopleID INT,  
    competitionID INT,  
    sportID INT,  
    oldrecord FLOAT,  
    newrecord FLOAT,  
    seton DATE  
);
```

- Function: BiggestRecordJump
 - Input: ID of sport
 - Output: The largest increase of that sport

```
RecordLog (  
    peopleID INT,  
    competitionID INT,  
    sportID INT,  
    oldrecord FLOAT,  
    newrecord FLOAT,  
    seton DATE  
);
```

Let's make a function (II)

- Assume RecordLog table:

```
RecordLog (  
    peopleID INT,  
    competitionID INT,  
    sportID INT,  
    oldrecord FLOAT,  
    newrecord FLOAT,  
    seton DATE  
);
```

- Function: BiggestRecordJump
 - Input: ID of sport
 - Output: The largest increase of that sport

```
DROP FUNCTION IF EXISTS BiggestRecordJump();
```

```
CREATE FUNCTION BiggestRecordJump(  
    IN sid INT  
)  
RETURNS FLOAT  
AS $$  
DECLARE r FLOAT; Variable to store value from query  
BEGIN  
    SELECT MAX(newrecord - oldrecord) INTO r  
    FROM RecordLog  
    WHERE sportID = sid  
    RETURN r;  
END;  
$$ LANGUAGE plpgsql;  
  
SELECT BiggestRecordJump(1);
```

Are Functions Faster?

- May be faster than executing from a client
 - Why?
- Code may be pre-compiled and optimized
 - Do not need to invoke optimizer again
 - May occur with well written queries using plan caching
- The code runs at the server
 - The server may be more efficient
 - No need to move data

Pros & Cons

- Code shared across ALL applications
- May be used for access control
- May give performance benefits
- Very system-specific
- Code maintenance requires care
 - Versioning is difficult

IDENTITY Columns

- Run 04-sports-schema.sql
- Try inserting a new person
- Due to the ID column not being an IDENTITY column it will fail inserting the row
- Quickfix: Add parameter pid in the NewPerson Function
- Best to avoid this by using `GENERATED ALWAYS AS IDENTITY` when creating the table
- Fixed tables with ID columns correctly set to `GENERATED ALWAYS AS IDENTITY` (run 04-sports-schema-fixed.sql)
- In case you want to insert a row with a specific ID, you can do so by `OVERRIDING SYSTEM VALUE`. See example in the file.

-- TODO

✓ Database Functions

- Database Triggers & Constraints
- Transactions
- SQL in Python



Triggers in SQL

- Automatically executed function in response to certain events on a table (or view) in a database
 - Events are typically INSERT, UPDATE, DELETE
 - Triggers can execute BEFORE or AFTER the event
- Useful for maintaining the integrity of the information on the Database.
 - For example, in a company database when a new record (representing a new worker) is added to the employees table, new records could also be created in the tables of taxes, vacations and salaries.
- Logging historical data, for example to keep track of employees' previous salaries.

Triggers in PostgreSQL

- Multiple triggers per table per event
 - INSERT / UPDATE / DELETE
 - Run in alphabetical order
 - Per row OR Per statement – focus here on per row
- New data is in the **NEW** record
 - For INSERT / UPDATE
 - Same schema as modified relation
 - Can refer to **NEW.ID**, **NEW.name**, ...
- Old data is in the **OLD** record
 - For UPDATE / DELETE
 - Same schema, refer to **OLD.ID**, **OLD.name**, ...
- The variable TG_OP says which operation it is

Trigger: Checking Values

- Use triggers to check value semantics
- A result in the Sports DB must not be negative

```
CREATE FUNCTION CheckResult()  
RETURNS TRIGGER  
AS $$ BEGIN  
    IF (NEW.result < 0.0) THEN  
        RAISE EXCEPTION  
        'CheckResult: Result must be a  
        positive'  
        USING ERRCODE = '45000';  
    END IF;  
    RETURN NEW;  
END; $$ LANGUAGE plpgsql;
```

Create
Trigger based
on function

```
CREATE TRIGGER CheckResult  
BEFORE INSERT OR UPDATE  
ON Results  
FOR EACH ROW EXECUTE PROCEDURE CheckResult();
```

Trigger: Ban Updates/Deletes


- Use triggers to not allow updating or deleting rows

```
CREATE FUNCTION BanChanges()  
RETURNS TRIGGER  
AS $$  
BEGIN  
    RAISE EXCEPTION  
        'BanChanges: Cannot change  
        results!'  
    USING ERRCODE = '45000';  
END;  
$$ LANGUAGE plpgsql;
```

```
CREATE TRIGGER BanChanges  
BEFORE UPDATE OR DELETE  
ON Results  
FOR EACH ROW EXECUTE PROCEDURE  
BanChanges();
```

Trigger: Update Records

- Use triggers to update the values after an event
- If a result in Sports DB is a new record, the row for that sport in the sports table need to be updated

```
CREATE FUNCTION UpdateRecord()  
RETURNS TRIGGER  
AS $$ BEGIN  
    IF NEW.result > (  
        SELECT s.record  
        FROM Sports s  
        WHERE s.id = NEW.sportID  
    ) THEN  
        UPDATE Sports  
        SET record = NEW.result  
        WHERE s.id = NEW.sportID;  
    END IF;  
    RETURN NEW;   
END; $$ LANGUAGE plpgsql;
```

Could also be **RETURN NULL** since result of AFTER trigger is ignored

```
CREATE TRIGGER UpdateRecord  
AFTER INSERT OR UPDATE  
ON Results  
FOR EACH ROW EXECUTE PROCEDURE UpdateRecord();
```

Trigger: Log Changes

```
CREATE TABLE RecordLog (  
    peopleID INT,  
    competitionID INT,  
    sportID INT,  
    oldrecord FLOAT,  
    newrecord FLOAT,  
    seton DATE  
    PRIMARY KEY  
    (peopleID, competitionID, sportID)  
    FOREIGN KEY  
    (peopleID, competitionID, sportID)  
    REFERENCES Results  
    (peopleID, competitionID, sportID)  
);
```

```
CREATE FUNCTION LogRecord()  
RETURNS TRIGGER  
DECLARE oldRecord FLOAT;  
AS $$ BEGIN  
    IF NEW.result > (  
        SELECT s.record  
        FROM Sports s  
        WHERE s.id = NEW.sportID  
    ) THEN  
        SELECT s.record INTO oldRecord  
        FROM Sports s  
        WHERE s.id = NEW.sportID;  
        INSERT INTO RecordLog  
        VALUES (NEW.peopleID, NEW.competitionID,  
            NEW.sportID, oldRecord,  
            NEW.result);  
    END IF;  
    RETURN NEW;  
END; $$ LANGUAGE plpgsql;  
  
CREATE TRIGGER LogRecord  
AFTER INSERT ON Results  
FOR EACH ROW EXECUTE PROCEDURE LogRecord();
```


Questions

- Why not create the LogRecord on Sports?
- What about Trigger order?
 - LogRecord vs UpdateRecord
 - Order is alphabetical = good in this case
 - Try with different order
 - May be better to merge similar triggers
- What about BEFORE or AFTER?
 - LogRecord and UpdateRecord are AFTER
 - CheckResult and BanUpdates are BEFORE
 - BEFORE and AFTER what? When does each apply? Why?

Merged Trigger (BanUpdates + Check Updates)

```
CREATE FUNCTION MergedTrigger()  
RETURNS TRIGGER  
AS $$ BEGIN  
    IF (TG_OP = 'DELETE' OR TG_OP = 'UPDATE')  
    THEN  
        RAISE EXCEPTION  
        'MergedTrigger: Cannot change results!'  
        USING ERRCODE = '45000';  
    END IF;  
    IF (NEW.result < 0.0) THEN  
        RAISE EXCEPTION  
        'CheckResult: Result must be a positive!'  
        USING ERRCODE = '45000';  
    END IF;  
END; $$ LANGUAGE plpgsql;
```

```
CREATE TRIGGER MergedTrigger  
BEFORE INSERT OR UPDATE OR DELETE  
ON Results  
FOR EACH ROW EXECUTE PROCEDURE  
MergedTrigger();  
  
INSERT INTO Results  
VALUES (1,1,3,-1.0);  
DELETE FROM Results WHERE sportID = 3;
```

BEFORE vs AFTER

1. Are you only checking the newly inserted/updated entry? -> **BEFORE (or AFTER)**
 - Checking happens earlier with BEFORE, so less work
 - But if other triggers might modify the values, then prefer AFTER
2. Are you inserting a row to another table with a foreign key constraint to the NEW record? -> **AFTER**
 - Otherwise, the NEW record is NOT in the database, so your insertion will fail!
3. Are you modifying the NEW record? -> **BEFORE**
 - Otherwise, the record is already in the database and will not be changed
4. Are you doing both 2 and 3? -> **BEFORE and AFTER**
 - Two different triggers!

Are Triggers Faster?

- May be faster than executing the code from a client
 - Why?
- Code may be pre-compiled and optimized
 - Do not need to invoke optimizer again
 - But this may happen with well written queries using plan caching
- The code runs at the server
 - The server may be more efficient
 - No need to move data to client

Pros & Cons

- Code often runs faster
 - No context switch
 - Compiled code
 - No data transfers
- Useful for security
 - Wrap data and functionality
 - Same access from all clients
- Same code for all applications
- Code is “hidden”
 - Only visible via system tables
 - Easily forgotten
 - Versioning is hard!
 - Schema may be edited using a GUI = no-no!
- Generally not portable

Exercises on SQL Programming

- Exercise 4 are out: Views, functions, triggers
 - Use test script to verify your work
 - Can run parts in pgAdmin – best via command prompt
 - You can extend the script with your own tests!
 - Also has a very nice database cleaning script

```
# once
psql -q Bills < bills-schema.sql

#repeatedly
psql Bills < universal-cleanup.sql
psql Bills < your-solution.sql
psql Bills < test-script.sql > output.txt 2>&1
less output.txt (linux / mac)
cat output.txt (windows) or open in an editor
```

NOTE: Windows uses '-f' instead of '<' and requires the DB name at the end of the command

NOTE 2: You may also be required to add user, -U <username>

Exercises on SQL Programming (Catalog)

- All information about the database is stored in tables!
 - This is the catalog!

<https://www.postgresql.org/docs/current/catalogs.html>

- See also: System information functions

<https://www.postgresql.org/docs/current/functions-info.html>

```
SELECT *  
  FROM information_schema.tables  
 WHERE table_schema = 'public';
```

```
SELECT current_database();
```

```
SELECT current_user;
```

```
SELECT lastval();
```

-- TODO

- ✓ Database Functions
- ✓ Database Triggers & Constraints
 - Transactions
 - SQL in Python





Transfer 50 DKK from Bob to Alice

Step 1: Get balance of Bob

```
SELECT Balance INTO x
FROM accounts
WHERE AccountName = 'Bob'
```

Account Name	Balance
Bob	200
Alice	0

Account Name	Balance
Bob	150
Alice	0

Step 2: Reduce balance of Bob

```
UPDATE accounts
SET Balance = x-50
WHERE AccountName = 'Bob'
```

WHAT IF SOMETHING CRASHES?

Step 3: Get balance of Alice

```
SELECT Balance INTO y
FROM accounts
WHERE AccountName = 'Alice'
```

Account Name	Balance
Bob	150
Alice	0

Account Name	Balance
Bob	150
Alice	50

Step 4: Increase balance of Alice

```
UPDATE accounts
SET Balance = y+50
WHERE AccountName = 'Alice'
```

What is a Transaction?

- A group of related operations to the database
- Wish for “all or nothing” execution
- Wish for isolation from other transactions

Basic SQL Syntax:

```
BEGIN;
```

```
COMMIT;
```

```
ROLLBACK;
```

Savepoints:

```
SAVEPOINT <name>;
```

```
ROLLBACK TO SAVEPOINT <name>;
```

Transactions in PostgreSQL

- By default, every statement is a transaction
 - To override this behaviour:
 - `BEGIN; ... COMMIT; / ROLLBACK;`
 - Some (DDL) statements implicitly `COMMIT` transactions
- Calling a function starts a transaction
 - Can assume that the function has transactional properties!
 - Errors abort the transaction, erase all previous operations!
- Errors inside functions:
 - Cannot simply say: `ROLLBACK`
 - Need to raise and handle exceptions!
 - See example in exercise code

Transactions and Testing

- Transactions are useful in test scripts
 - Multiple examples in Exercise 4
- Simple pattern to test changing the database
 - ... without actually changing the database

```
BEGIN
```

```
-- Make changes
```

```
-- Run test queries
```

```
ROLLBACK
```

-- TODO

- ✓ Database Functions
- ✓ Database Triggers & Constraints
- ✓ Transactions
 - SQL in Python


Problem: Vendor Lock-in

- They're all compatible with standard SQL...
- ...and mutually incompatible due to vendor-specific functionality!



Open Database Connectivity (ODBC)

- In-between translation layer
- Vendors not so happy as lock-in is profitable!



Python and PostgreSQL

- pyodbc is a package that implements ODBC and can connect to any databases
- The psycopg package is a database adapter specifically for PostgreSQL
 - Reliable, documented, references
 - To install: `python -m pip install psycopg` or `pip install psycopg`
 - `import psycopg`
- A simple database API - Similar to ODBC
 - `connect()`: Establish connection to specified database
 - `cursor()`: Object to manage context of an SQL operation
 - `execute()`: Execute SQL queries
 - `fetchone()`, `fetchmany()`, `fetchall()`: Get results from queries
 - Much more, but these are the most used functions

Connecting to a Database

```
import psycopg2 as pg
from psycopg2.rows import dict_row ← Optional: Rows are returned as dict, better readability


conn_string = "host=localhost dbname=Sports user=postgres password=*****"

# Connection option one
conn = pg.connect(conn_string, autocommit=True, row_factory=dict_row)
# Execute Queries using Cursors
conn.close()

# Connection option two
with pg.connect(conn_string, autocommit=True, row_factory=dict_row) as conn:
    # Execute Queries using Cursors
# Closes connection
```

Tips

```
try:
    # Connection option one
    conn = pg.connect(conn_string, autocommit=True, row_factory=dict_row)
    # Execute Queries using Cursors
except Exception as error:
    print(error)
finally:
    conn.close()
```



```
try:
    # Connection option two
    with pg.connect(conn_string, autocommit=True, row_factory=dict_row) as conn:
        # Execute Queries using Cursors
    # Closes connection
except Exception as error:
    print(error)
```

Querying the DB

Get all athletes

```
athletes: list[dict] = conn.execute("SELECT * FROM Athletes").fetchall()
```

Creates a cursor object and
executes the query

Get five athletes

```
athletes: list[dict] = conn.execute("SELECT * FROM Athletes").fetchmany(5)
```

What if I want more items later
from the same query?

Get all athletes

```
athletes: list[dict] = conn.execute("SELECT * FROM Athletes").fetchone()
```

Cursors

```
# Get query cursor
athletes_cursor = conn.execute("SELECT * FROM Athletes")

# Get five athletes
athletes: list[dict] = athletes_cursor.fetchmany(5)

# Get sixth athlete
athlete_six = athletes_cursor.fetchone()

# Get the next 10 athletes
athlete_next_ten = athletes_cursor.fetchmany(10)
```

Cursors

```
# Get query cursor
```

```
athletes_cursor = conn.execute("SELECT * FROM Athletes")
```

```
# Get five athletes
```

```
athletes: list[dict] = athletes_cursor.fetchmany(5)
```

```
sports_cursor = conn.execute("SELECT * FROM Sports")
```

```
# Get sixth athlete
```

```
athlete_six = athletes_cursor.fetchone()
```


```
sports: list[dict] = sports_cursor.fetchall()
```

```
# Get the next 10 athletes
```

```
athlete_next_ten = athletes_cursor.fetchmany(10)
```

```
athletes_cursor.close()
```

```
sports_cursor.close()
```



Advantage: Multiple cursors
active at the same time, which
can fetch or execute queries

Remember to close the cursors (context block, finally block)

Creating a Table

```
with pg.connect(conn_string, autocommit=True, row_factory=dict_row) as conn:
    conn.execute("DROP TABLE IF EXISTS Quotes")
    conn.execute(
        """
        CREATE TABLE Quotes (
            x INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
            y VARCHAR(250) NOT NULL,
        )
        """
    )
```

Better than SERIAL and follows SQL Standard.
Still a sequence under the hood.
[PostgreSQL Identity Columns](#)

INSERT

```
conn = pg.connect(conn_string, autocommit=True, row_factory=dict_row)
with conn:
    conn.execute("INSERT INTO Quotes (y) VALUES ('I am Vengeance')")

nope = conn.execute("SELECT 42 as universe").fetchone()
```

← Won't run since the connection is closed

INSERT

```
conn = pg.connect(conn_string, autocommit=True, row_factory=dict_row)
with conn:
```

```
    conn.execute("INSERT INTO Quotes (y) VALUES ('I am Vengeance')")
```

```
nope = conn.execute("SELECT 42 as universe").fetchone()
```

← Won't run since the connection is closed

```
with conn.cursor() as cur:
```

```
    cur.execute("INSERT INTO Quotes (y) VALUES ('I am the Night')")
```

```
yes = conn.execute("SELECT 42 as universe").fetchone()
```

← Will run since only the cursor closed

```
conn.close()
```

Querying with variables

```
in_text = 'I am Batman'  
sel_text = "'%Batman%'"
```

```
with pg.connect(conn_string, autocommit=True, row_factory=dict_row) as conn:  
    conn.execute("INSERT INTO Quotes (y) VALUES (%s)", [in_text])  
    conn.execute("SELECT * FROM Quotes WHERE y LIKE %s" % sel_text)
```

Question: Is this good or bad?


SELECT Query is Prone to SQL Injections!

SQL Injection Example

```
in_text = 'I am Batman'  
sel_text = "'%Batman%'; DELETE FROM Quotes WHERE 1=1; INSERT INTO Quotes (y)  
VALUES 'Joker was here';--"
```

```
with pg.connect(conn_string, autocommit=True, row_factory=dict_row) as conn:  
    conn.execute("INSERT INTO Quotes (y) VALUES (%s)", [in_text])  
    conn.execute("SELECT * FROM Quotes WHERE y LIKE %s" % sel_text)
```

Even relevant in the future:



Star Trek Discovery Season 2
Episode 8

Year 2255+

Prepared Statements

- Separate the query string and variables
- Prepared Statements
 - Pre-compiled SQL statement at DBMS
 - Client supplied data is treated as content of a parameter and not an SQL statement

```
PREPARE my_query AS  
SELECT *  
FROM Quotes  
WHERE y = "$1"
```

```
EXECUTE my_query ("Probe used SQL  
Injection... it's not very effective")
```

Prepared Statements in Python

```
in_text = 'I am Batman'  
sel_text = "'%Batman%'; DELETE FROM Quotes WHERE 1=1; INSERT INTO Quotes (y)  
VALUES 'Joker was here';--"
```

```
with pg.connect(conn_string, autocommit=True, row_factory=dict_row) as conn:  
    conn.execute("INSERT INTO Quotes (y) VALUES (%s)", [text])  
    conn.execute("SELECT * FROM Quotes WHERE y LIKE %s", [sel_text])
```



Not safe against all cases!

- Some elements can't be prepared
- Most prominent is the ORDER BY clause
- Check the input in python before using the variable

```
sort_key = "name"

if sort_key in columns:

    cur = conn.execute(
        f"""
        SELECT *
        FROM Athletes
        ORDER BY {sort_key}
        """)

    athletes = cur.fetchall()
```


Autocommit

- If not specified in `connect()` autocommit is OFF by default
- Using `with` connection context commits at the end of the block

```
with pg.connect(conn_string, row_factory=dict_row) as conn:
    conn.execute("INSERT INTO Quotes (y) VALUES ('I am the one who knocks')")
    # BEGIN;
    # INSERT...;
    # COMMIT; (default) / ROLLBACK; (if error is raised)
```

- Good for a small number of queries
- But closes connection after, so this may not always be desired

Autocommit (Cursors)

```
conn = pg.connect(conn_string, row_factory=dict_row)
with conn.cursor() as cur:
    tests = cur.execute("SELECT * FROM Quotes")
    # BEGIN;
    # SELECT...;
# Not committed, manually call conn.commit() inside or after the block
conn.commit()
```

- Even a SELECT statement begins a transaction
 - Easy to forget, making it more dangerous to not set autocommit on

Transactions with Autocommit ON

```
conn = pg.connect(conn_string, autocommit=True, row_factory=dict_row)

conn.execute("INSERT INTO Quotes (y) VALUES ('I am inevitable')")

# BEGIN; SELECT...; COMMIT/ROLLBACK;

with conn.transaction():
    # BEGIN;
    with conn.cursor() as cur:
        cur.execute("INSERT INTO Quotes (y) VALUES ('I am Groot')")
        cur.execute("INSERT INTO Quotes (y) VALUES ('I am Iron Man!')")
    # COMMIT; (default) / ROLLBACK; (if error is raised)
```

-- TODO -> DONE

- ✓ Database Functions
- ✓ Database Triggers & Constraints
- ✓ Transactions
- ✓ SQL in Python

Takeaways

Functions

Set of database operations, performance benefits

Triggers

Execute a function in response to a database event

ODBC (Python `psycopg`, Java `JDBC`)

API for database operations

Transactions

Atomicity, Consistency, Isolation, Durability



Next Time in IDBS...

Introduction to Database Systems

IDBS - Fall 2024

Lecture 5 - Designing Databases

ER Diagrams

Translation to SQL DDL

Readings: PDBM 3.0-3.3, 6.3-6.4

Omar Shahbaz Khan

Extra bits

Roles in Databases

- Multiple users interacting with the database
 - Should everyone use the postgres role?
 - Roles are the same as Users
- Access Control
 - We can define roles with different privileges
 - Avoid accidents or fatal errors (database related and business related)

Roles in SQL

- CREATE

```
CREATE ROLE <role> WITH LOGIN ENCRYPTED PASSWORD <password>;  
CREATE ROLE viewer WITH LOGIN ENCRYPTED PASSWORD '1337';
```

- GRANT

```
GRANT <privilege> ON <schema_elements> TO <role>;  
GRANT SELECT ON Sports, Results, Competitions, Athletes,  
Gender TO viewer;
```

- REVOKE

```
REVOKE SELECT ON Sports, Results, Competitions, Athletes,  
Gender FROM viewer;
```

- DROP OWNED BY

```
DROP OWNED BY viewer;
```

- DROP

```
DROP ROLE viewer;
```

Using Roles

- To use roles in Python and in general we use the SQL statement SET ROLE
- SET ROLE <role>
- SET LOCAL ROLE <role>
- Best to execute in a transaction so we do not accidentally allow others to execute as the set role
- Check the add functions in the exercise

Roles and Functions I

- Creating a function in PostgreSQL grants execute privileges on public
 - Many roles could have access to public!
 - Solution: Revoke access from public and then grant access to specific roles
 - **REVOKE ALL ON FUNCTION** <function> **FROM** public;
 - **GRANT EXECUTE ON FUNCTION** <function> **TO** <role>;

Roles and Functions II

- By default a function is executed with the privileges of the calling role
 - A function accessing a schema element which the calling role does not have privileges to, an `InsufficientPrivileges` error is raised
- Resolution: **SECURITY DEFINER**
 - Add this after **LANGUAGE plpgsql**
 - The function is run with the privileges of the user that created it
 - Be careful!

A C I D

Transaction Properties

Recovery

Constraints, Triggers

Concurrency
Control

Recovery

A TOMICITY

A Transaction is
“one operation”

CONSISTENCY

Each Transaction
moves the DB from one
consistent state to
another

ISOLATION

Each Transaction
is alone in the
world

DURABILITY

Persistence of
successful
transactions even
through system failure

Transaction Implementation Methods (Lecture 10)

- Consistency
 - PRIMARY and FOREIGN KEY = limited DBMS support
- Isolation
 - Historically locking = (strict/rigorous) two-phase locking
 - Recently multi-version concurrency control
- Atomicity / Durability
 - Logging all changes to disk
 - Write Ahead Logging protocol (WAL)