

# Introduction to Database Systems

## IDBS - Fall 2024

### Lecture 5 - Designing Databases

ER Diagrams

Translation to SQL DDL

---

Readings: PDBM 3.0-3.3, 6.3-6.4

Omar Shahbaz Khan

# General Info

## **EXERCISE 5: ER Design and Implementation**

- Large and comprehensive
- More than 2hrs, but use it to practice

## **HOMEWORK 2 - OUT NOW!**

- DEADLINE: 14. October 2024 23:59

# -- TODO

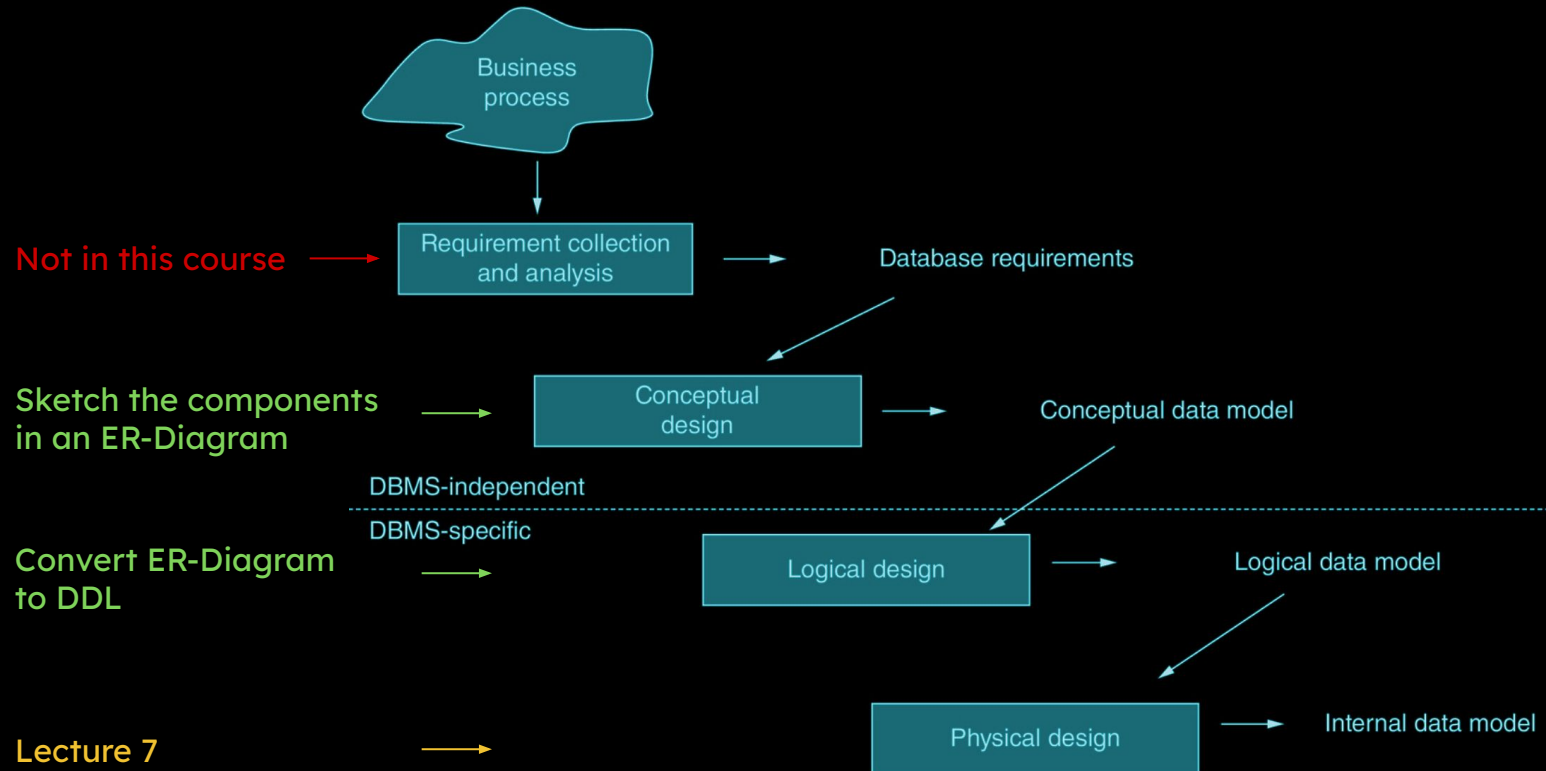
- Conceptual Data Modeling (ER Diagram)
  - Entities and Attributes
  - Relationships
    - Cardinalities
    - Partial Relationship Keys
  - Weak Entities
  - Aggregation
  - Generalization/Specialization
  - Categorization
- Translation to SQL DDL

# Conceptual Data Modeling

# Why do we need a Conceptual Model?

- The “higher-ups” know they want a database...  
... but not what should be in it!
- Need an effective method to develop the schema and document its structure

# Design Process



# ER: Entity-Relationships

- Conceptual Data Modeling technique was defined by Peter Chen (1976)
- ER = modeling concepts + visual representation
  - ER/EER notation is not standardized
  - Every textbook/company/tool has its own visual representation
  - Core concepts are universally accepted
  - EER vs ER rarely distinguished -> we just call it ER!
- UML can be used as design notation
  - Slight differences – NO UML IN THIS COURSE
- Focus on ER notation from the book
  - ... plus, some minor extensions – made clear in the lectures
  - In exercises, project, exam: Use the notation in book, lectures


# Entity and Attribute Types

## Entity Type:

- Set of similar “things”
  - Ex: Movie, Supplier
- Entity = Instance
  - Ex: Interstellar, 2014

## Attribute Type:

- Describes one aspect of an entity set
  - Ex: name, year, address

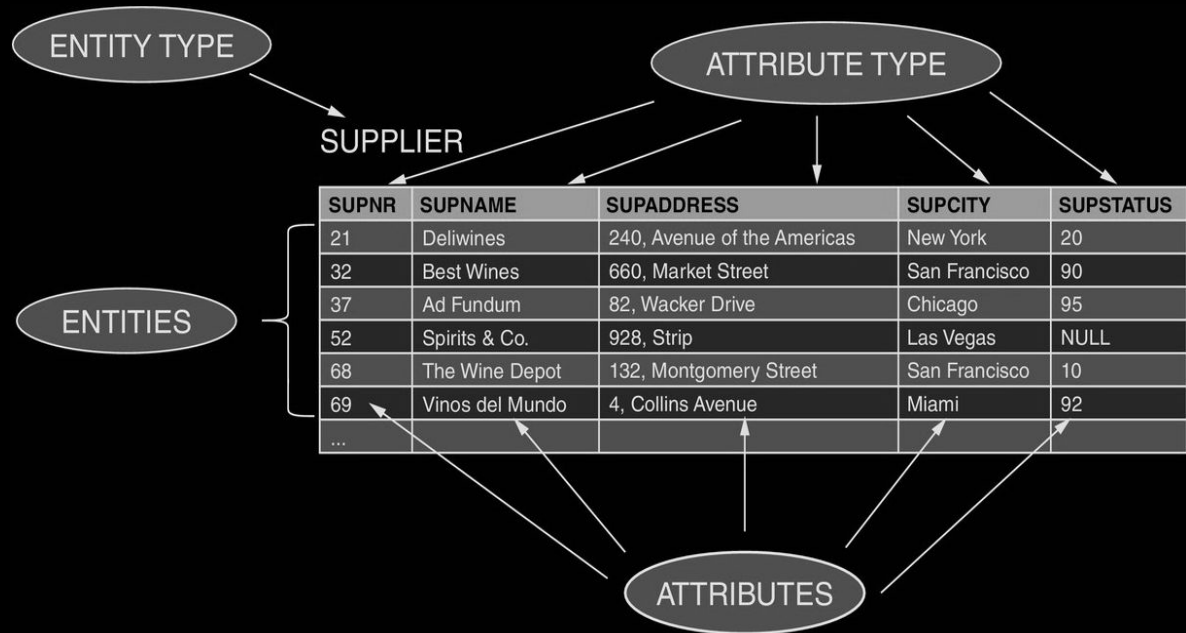



Is there something suspicious about this entity?

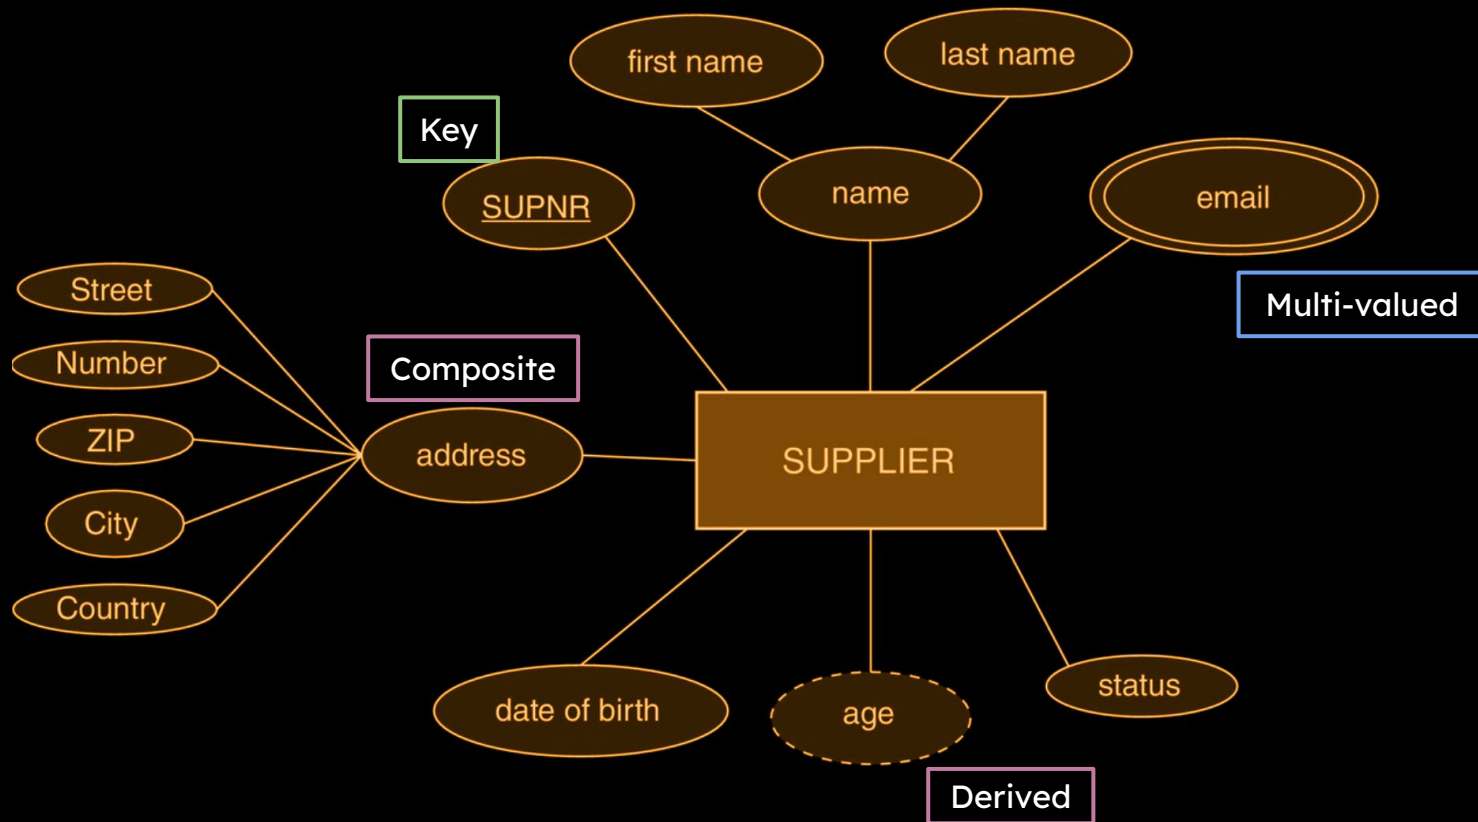


# Entity and Attribute Types

```
CREATE TABLE Supplier (
  SUPNR INT PRIMARY KEY,
  SUPNAME VARCHAR NOT NULL,
  ...
);
```



# More Attribute Types





# Keys in ER-Diagrams and SQL DDL

- Underlined key = PRIMARY KEY
- ER diagrams cannot show secondary keys
  - They must be noted somewhere else!
  - They must still be UNIQUE in the SQL table!

```
CREATE TABLE Supplier (
    SUPNR INT PRIMARY KEY,
    ...
);
```

```
CREATE TABLE Supplier (
    SUPNR INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
    ...
);
```

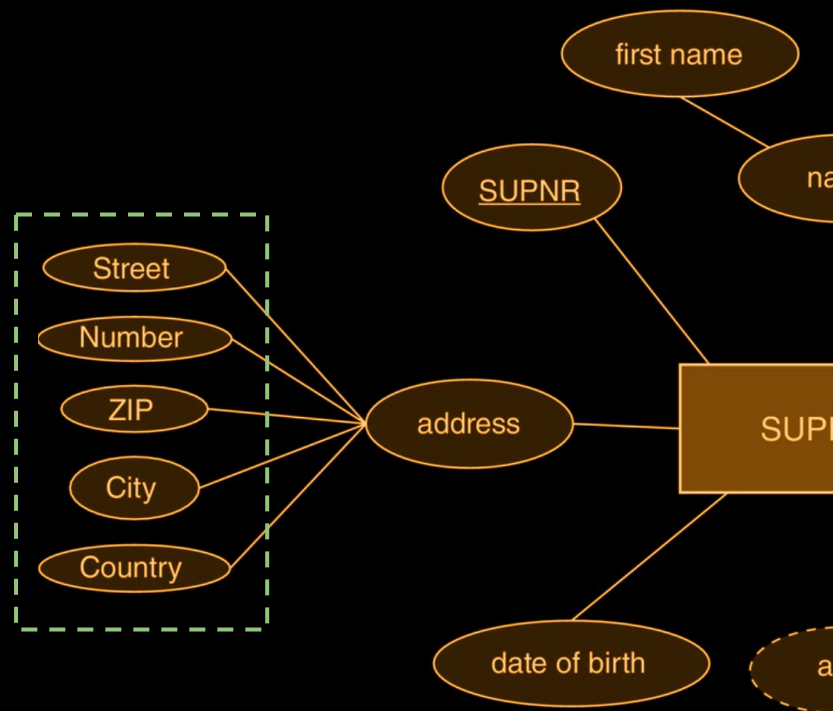

In case you need an  
incremental sequence





# Composite Attributes in SQL DDL

- Simply add the detailed attributes:

```
CREATE TABLE Supplier (
  SUPNR INT PRIMARY KEY,
  ...
  Street VARCHAR NOT NULL,
  Number INTEGER NOT NULL,
  ZIP INTEGER NOT NULL,
  City VARCHAR NOT NULL,
  Country VARCHAR NOT NULL,
  ...
);
```



# Multi-Valued Attributes in SQL DDL






```
CREATE TABLE Emails (  
    Email VARCHAR,  
    SupNR INTEGER REFERENCES Supplier(SupNR),  
    PRIMARY KEY (Email, SupNR)  
);
```

- Could add `ON DELETE CASCADE` to SupNR
- Generally a good idea to add to foreign keys, if the related rows should not exist independent of that key

# Derived Attributes in SQL DDL

- **Not discussed in the PDBM book!**
- Option 1: Create an attribute and maintain it
  - E.g. with a trigger, or regular update processes
- Option 2: Create a view that computes it
- Neither is very good!
- Sometimes there is a second kind of inter-attribute relationship
  - Here: ZIP  $\rightarrow$  City
  - Called functional dependencies (FDs)
  - ER diagrams may miss such relationships!
  - We fix this with normalization (Lecture 6)



# The CHECK Statement

- We can add constraints that check the value of a field directly in the DDL
- These can be on the column or on the table
  - Postgresql does not care but other database system may

<https://www.postgresqltutorial.com/postgresql-tutorial/postgresql-check-constraint/>

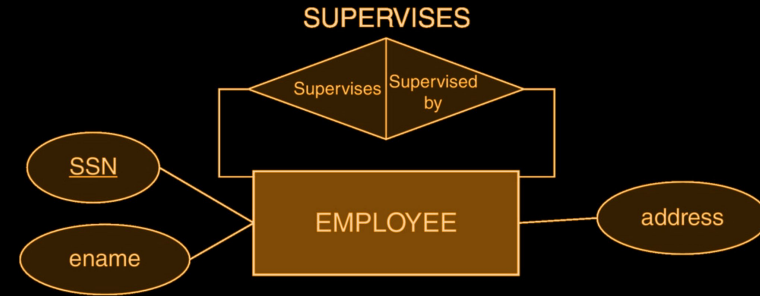

```
CREATE TABLE Supplier (  
    Id INTEGER PRIMARY KEY,  
    ...  
    status CHAR(10) NOT NULL  
    CHECK (status IN ('Active', 'Inactive')),  
    ...  
);  
  
CREATE TABLE Product (  
    Id INTEGER PRIMARY KEY,  
    ...  
    price FLOAT NOT NULL,  
    ...  
    CHECK (price > 0),  
    ...  
);
```

# Relationships



# Relationship Types

- Relate two or more entities (with roles)
- Ex: John majors in Computer Science
- Roles may be omitted when obvious



```
CREATE TABLE Supplies (
  SupNR INT REFERENCES Supplier,
  ProdNR INT REFERENCES Product,
  PRIMARY KEY (SupNR, ProdNR)
);
```

# ER-Diagram Exercise


## DB Requirements:

- A drinker has a (unique) name, address, and phone.
- A coffeehouse has a (unique) name, address, and license.
- Store which drinkers frequent which coffeehouses.

# ER-Diagram Exercise


## DB Requirements:

- A drinker has a (unique) name, address, and phone.
- A coffeehouse has a (unique) name, address, and license.
- Store which drinkers frequent which coffeehouses.




# Relationship Attribute Types

- Relationships can also have attribute types



# Example: Coffee DB

- What if...
  - price is an attribute of Coffee?
  - price is an attribute of Coffeehouse?



Coffees (name, manufacturer)

Coffeehouses (name, address, license)


Drinkers (name, address, phone)

Likes (drinker, coffee)

Sells (coffeehouses, coffee, price)

Frequents (drinker, coffeehouse)


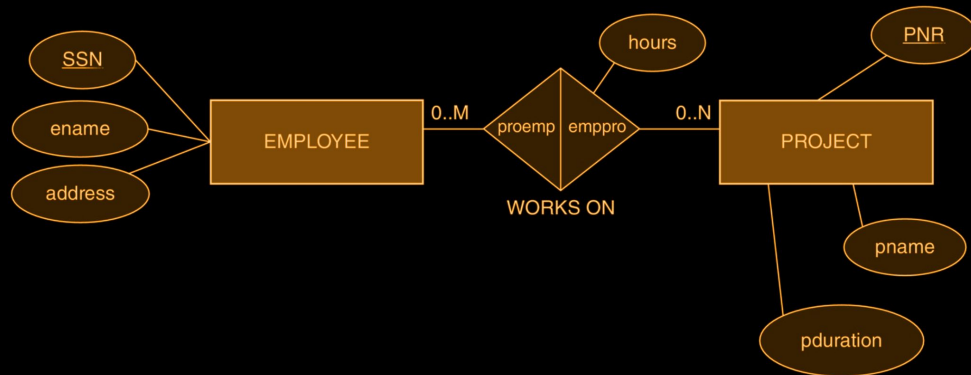
# Basic Relationship Table in SQL DDL




```
CREATE TABLE WorksIn (  
    ProfID INT,                -- role (key of Professor)  
    DeptID CHAR(4),            -- role (key of Department)  
    Since DATE NOT NULL,       -- attribute  
    Status CHAR(10) NOT NULL,  -- attribute  
    PRIMARY KEY (ProfID, DeptID),  
    FOREIGN KEY (ProfID) REFERENCES Professor (ID),  
    FOREIGN KEY (DeptID) REFERENCES Department (ID)  
)
```

# Cardinalities

- Relationships always have cardinalities
  - Minimum: 0 or 1
  - Maximum: 1 or N / M / P / \* / ...
- Read:
  - Entity (ignore) Relationship Cardinality Entity
  - Employee can work on zero to many projects
- Do cardinalities impact the resulting table structure?



# Maximum 1



```
CREATE TABLE WorksIn (  
    ProfID INT,                -- role (key of Professor)  
    DeptID CHAR(4) NOT NULL,,  -- role (key of Department)  
    Since DATE NOT NULL,      -- attribute  
    Status CHAR(10) NOT NULL, -- attribute  
    PRIMARY KEY (ProfID),     -- each professor only once  
    FOREIGN KEY (ProfID) REFERENCES Professor (ID),  
    FOREIGN KEY (DeptID) REFERENCES Department (ID)  
);
```




# Maximum 1 - No Attributes



~~CREATE TABLE WorksIn (...)~~

```
CREATE TABLE Professor (  
  Id INT PRIMARY KEY,  
  ...  
  deptId INT, -- foreign key  
  FOREIGN KEY (deptId) REFERENCES Department(Id)  
)
```


# Exactly 1



~~CREATE TABLE WorksIn (...)~~

```
CREATE TABLE Professor (  
  Id INT PRIMARY KEY,  
  ...  
  deptId INT NOT NULL, -- foreign key  
  Since DATE NOT NULL, -- attribute  
  Status CHAR(10) NOT NULL, -- attribute  
  FOREIGN KEY (deptId) REFERENCES Department(Id)  
)
```

# Maximum 1 - Revisited




Could we do the same for 0..1?

- Yes, but the three attributes must be able to be NULL
- **We only do this if the relationship has no attributes**

```
CREATE TABLE Professor (  
  Id INT PRIMARY KEY,  
  ...  
  deptId INTEGER, foreign key  
  Since DATE, attribute  
  Status CHAR(10), attribute  
  FOREIGN KEY (deptId) REFERENCES Department(Id)  
)
```


# Maximum 1 - Both directions



```
CREATE TABLE WorksIn (  
    ProfID INT,                -- role (key of Professor)  
    DeptID CHAR(4) NOT NULL,   -- role (key of Department)  
    Since DATE NOT NULL,      -- attribute  
    Status CHAR(10) NOT NULL, -- attribute  
    PRIMARY KEY (ProfID),      -- each professor only once  
    UNIQUE (DeptID),           -- each department only once  
    FOREIGN KEY (ProfID) REFERENCES Professor(ID),  
    FOREIGN KEY (DeptID) REFERENCES Department(ID)
```

)

# Exactly 1 - Both direction



```
CREATE TABLE Professor (  
    Id INT PRIMARY KEY,  
    ...  
    DeptId INT NOT NULL,  
    Since DATE NOT NULL,  
    Status CHAR(10) NOT NULL,  
    FOREIGN KEY (DeptId)  
    REFERENCES Department(Id)  
)
```

```
CREATE TABLE Department (  
    Id INT PRIMARY KEY,  
    ...  
    ProfId INT NOT NULL,  
    FOREIGN KEY (ProfId)  
    REFERENCES Professor(Id)  
)
```

## Exactly 1 - Both direction

- Can we use FK on both sides?
  - Yes... but it is neither easy nor portable
- Think about inserting the first prof and dept
  - Which comes first... Chicken or the egg?
- Alternative 1: Use Deferred FK or Trigger
  - Runs at the end of a transaction - many systems support neither!

```
CREATE TABLE Professor (  
    Id INT PRIMARY KEY,  
    ...  
    DeptId INT NOT NULL,  
    Since DATE NOT NULL,  
    Status CHAR(10) NOT NULL,  
    FOREIGN KEY (DeptId)  
    REFERENCES Department(Id)
```

```
)
```


```
CREATE TABLE Department (  
    Id INT PRIMARY KEY,  
    ...  
    ProfId INT NOT NULL,  
    FOREIGN KEY (ProfId)  
    REFERENCES Professor(Id)  
)
```

# Exactly 1 - Both directions

- Alternative 2: Merge Tables
  - May work well in some cases
  - Depends on entities
- Alternative 3: Pick one FK direction
  - Write down the other requirement
  - Do the best in software with the other direction

```
CREATE TABLE Professor (  
    Id INT PRIMARY KEY,  
    ...  
    deptId INT NOT NULL,  
    deptName VARCHAR NOT NULL,  
    ...  
    Since DATE NOT NULL,  
    Status CHAR(10) NOT NULL  
)
```

# Minimum 1 - Maximum N



- What about 1..N or 1..M cardinalities?
  - Or when requirements demand particular fixed numbers?
- No support in SQL DDL
  - Could use a trigger on Professor/WorksIn
- Normally:
  - Write down the requirement
  - Do the best in software





# Exercise

- Draw this schema as ER diagram...

BUT:


- Use IDs
- Assume each coffeehouse sells exactly one coffee

Coffees(name, manf)

Coffeehouses(name, addr, license)


Sells(coffeehouse, coffee, price)

- Write SQL DDL to create the tables
  - How many tables?




## Exercise Answer DDL

```
CREATE TABLE Coffee (  
  ID INT PRIMARY KEY,  
  name VARCHAR NOT NULL,  
  manf VARCHAR NOT NULL  
);  
  
CREATE TABLE Coffeehouse (  
  ID INT PRIMARY KEY,  
  name VARCHAR NOT NULL,  
  addr VARCHAR NOT NULL,  
  lic VARCHAR NOT NULL,  
  coffeeID INT NOT NULL REFERENCES Coffee,  
  price INT NOT NULL  
);
```



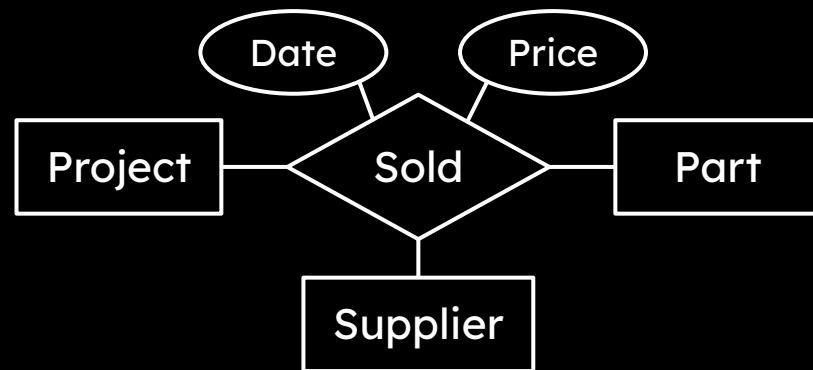
# Ternary Relationships (and Beyond)




- An Instructor can offer many Courses during a Semester
- A Course offered in a given Semester must have at least one Instructor involved
- An Instructor can offer a given Course for many Semesters

## Example: DLL of Ternary Relationship

```
CREATE TABLE Sold (  
  ProjID INT, -- role  
  SupplierID INT, -- role  
  PartNumber INT, -- role  
  Date DATE NOT NULL, -- attribute  
  Price FLOAT NOT NULL -- attribute  
 | PRIMARY KEY (ProjID, SupplierID, PartNumber),  
  FOREIGN KEY (ProjID) REFERENCES Project (ID),  
  FOREIGN KEY (SupplierID) REFERENCES Supplier (ID),  
  FOREIGN KEY (PartNumber) REFERENCES Part (Number)  
)
```



## Extension: Partial Relationship Keys




**PRIMARY KEY** (ProjID, SupplierID, PartNumber)

- What does the key of the relationship table mean?
- What if the part should be sold many times?
- The book: No discussion! :(
- Our notation: *Partial relationship keys*
  - Attribute is underlined

## Extension: Partial Relationship Keys in SQL DDL

```
CREATE TABLE Sold (  
  ProjID INT, -- role  
  SupplierID INT, -- role  
  PartNumber INT, -- role  
  Date DATE NOT NULL, -- attribute  
  Price FLOAT NOT NULL -- attribute  
  PRIMARY KEY (ProjID, SupplierID, PartNumber, Date),  
  FOREIGN KEY (ProjID) REFERENCES Project (ID),  
  FOREIGN KEY (SupplierID) REFERENCES Supplier (ID),  
  FOREIGN KEY (PartNumber) REFERENCES Part (Number)  
)
```



# -- TODO

- Conceptual Data Modeling (ER Diagram)


- ✓ Entities and Attributes
- ✓ Relationships
  - ✓ Cardinalities
  - ✓ Partial Relationship Keys
- Weak Entities
- Aggregation
- Generalization/Specialization
- Categorization

- Translation to SQL DDL



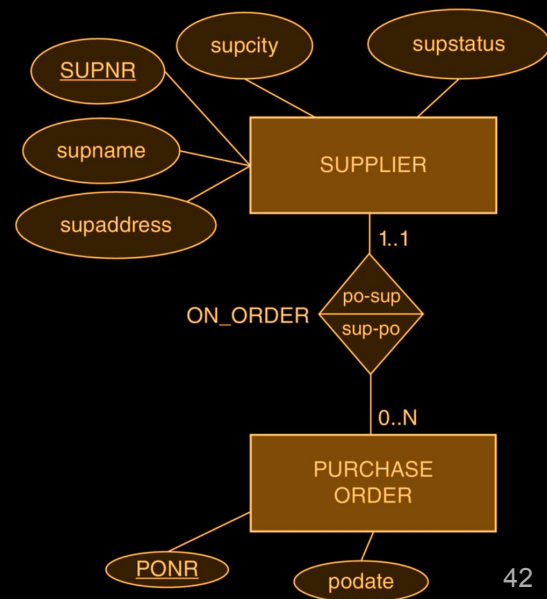
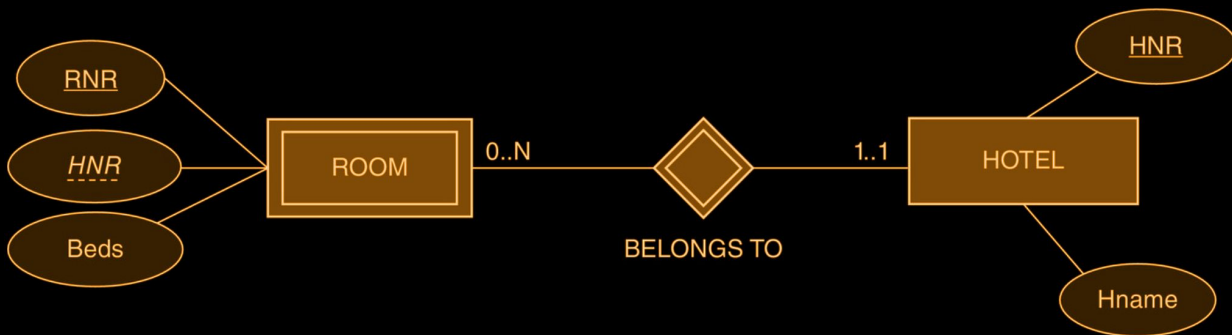
# Weak Entity Types

- Weak entities belong to another entity
  - They do not have a proper key, but include “parent” key
  - They have a 1 .. 1 participation in the relationship
  - If “parent” is deleted, so is the “child”
- Representation
  - Double outlines (entity, relationship)
  - Parent key is underlined with dashes



# Weak Entity Types vs 1..1 Relationship Types

- Weak entities have a 1..1 relationship type
  - Some 1..1 relationship types represent weak entities
  - Most 1..1 relationship types do not represent weak entities
- Main difference is presence of a natural key
  - Weak entity:  
Only unique within the parent entity!




## Practice Weak Entities

- Each Volleyball Match consists of Sets, each set has a number, home score and away score

## Practice Weak Entities


- Each volleyball match consists of sets, each with set number, home score and away score



# Weak Entities in SQL DDL


- Create a new table referencing the entity table
  - Primary key is parent key + partial key
  - Very similar to multi-valued attributes

```
CREATE TABLE Room (  
    RNR INT, -- Should not be a sequence!  
    HNR INT, -- Must be a FOREIGN KEY!  
    Beds INT NOT NULL,  
    PRIMARY KEY (HNR, RNR),  
    FOREIGN KEY REFERENCES Hotel (HNR),  
);
```

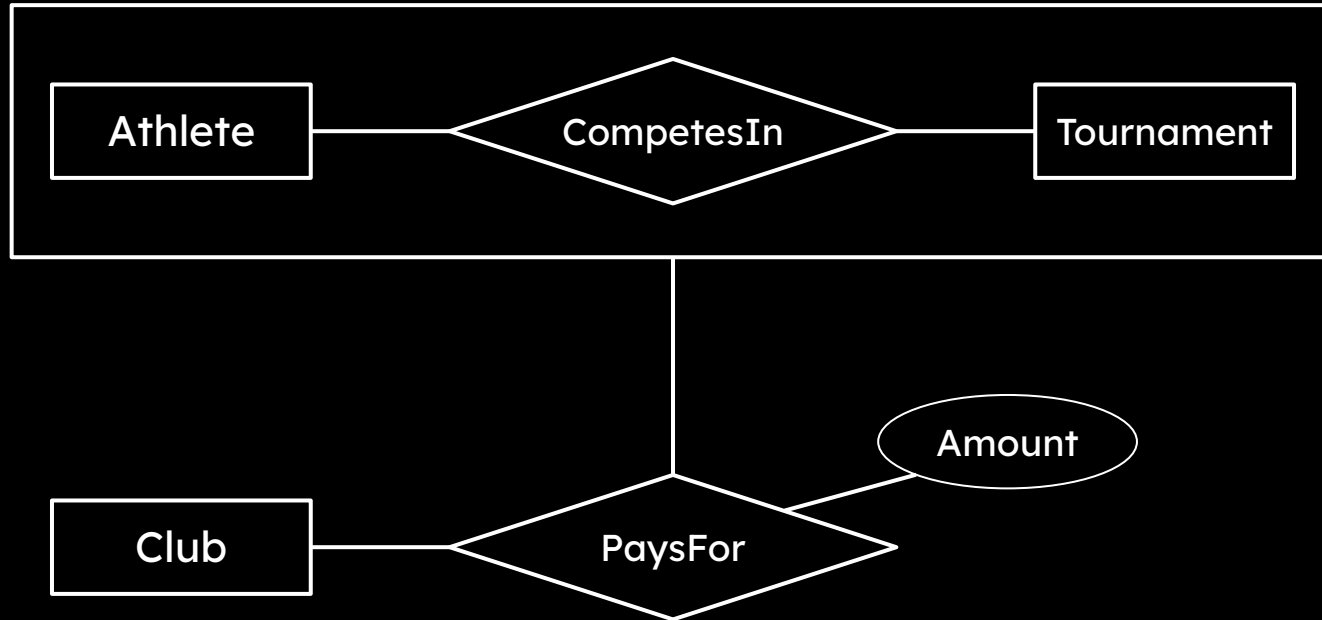


# Aggregation


- Sometimes we need a relationship to a relationship
- Aggregation allows us to “convert” relationship types to entity types!
- Typical use: Monitoring, payments, contracts



# Relationship → Entity




# Relationship → Entity






# Translation to SQL DDL

- Main observation: It is a relationship!
  - Use same method as translating relationship
  - Treat the aggregation table as the entity



# Translation to SQL DDL: 0..N

- Option 1: Use existing relationship key





```
CREATE TABLE CompetesIn (  
    AID INT REFERENCES Athlete,  
    TID INT REFERENCES Tournaments,  
    PRIMARY KEY (AID, TID)  
);
```

```
CREATE TABLE PaysFor (  
    AID INT,  
    TID INT,  
    CID INT REFERENCES Clubs,  
    amount INT NOT NULL,  
    FOREIGN KEY (AID, TID)  
    REFERENCES CompetesIn (AID, TID),  
    PRIMARY KEY (AID, TID, CID)  
);
```

# Entities and Aggregation: Impact on Tables


- A common error is to use FKs to entity tables



```
CREATE TABLE PaysFor (
  AID INT REFERENCES Athlete,
  TID INT REFERENCES Tournament,
  CID INT REFERENCES Clubs,
  amount INT NOT NULL,
  FOREIGN KEY (AID, TID)
    REFERENCES CompetesIn (AID, TID),
  PRIMARY KEY (AID, TID, CID),
);
```

# Translation to SQL DDL: 0..N

- Option 2: Create a new relationship key
  - Common error is to forget the existing key




```
CREATE TABLE CompetesIn (  
  CIID INT ALWAYS GENERATED AS IDENTITY  
    PRIMARY KEY,  
  AID INT REFERENCES NOT NULL Athletes,  
  TID INT REFERENCES  
    NOT NULL Tournament,  
  UNIQUE (AID, TID)  
);
```

```
CREATE TABLE PaysFor (  
  CIID INT REFERENCES CompetesIn,  
  CID INT REFERENCES Clubs,  
  amount INTEGER NOT NULL,  
  PRIMARY KEY (CIID, CID)  
);
```

# Translation to SQL DDL: 0..1

- An Athlete competing in a Tournament can receive payment from maximum one Club
  - Change the Primary Key




```
CREATE TABLE CompetesIn (  
    AID INT REFERENCES Athletes,  
    TID INT REFERENCES Tournament,  
    PRIMARY KEY (AID, TID)  
);
```

```
CREATE TABLE PaysFor (  
    AID INT,  
    TID INT,  
    CID INT NOT NULL REFERENCES Clubs,  
    amount INT NOT NULL,  
    FOREIGN KEY (AID, TID)  
    REFERENCES CompetesIn (AID, TID),  
    PRIMARY KEY (AID, TID)  
);
```

# Translation to SQL DDL: 1..1

- Change the relationship table
- This is the only case covered in the book!
- Here: No PaysFor table!



```
CREATE TABLE CompetesIn (  
    AID INT REFERENCES Athletes,  
    TID INT REFERENCES Tournament,  
    CID INT NOT NULL REFERENCES Clubs,  
    amount INT NOT NULL  
    PRIMARY KEY (AID, TID)  
);
```

# -- TODO -> DONE



## ✓ Conceptual Data Modeling (ER Diagram)

- ✓ Entities and Attributes
- ✓ Relationships
  - ✓ Cardinalities
  - ✓ Partial Relationship Keys
- ✓ Weak Entities
- ✓ Aggregation
  - Generalization/Specialization
  - Categorization

## ● Translation to SQL DDL

# Generalization/Specialization

- Like inheritance in Object-Oriented Programming (Java/C#/etc.)
- Partial vs Total = p/t on the line
- Overlapping vs Disjoint = o/d in the circle
- Please don't use colour in homeworks/exam
- Arcs matter → need to draw them (in the exam)!





# Specialization in SQL DDL

- One table for super-type, one per sub-type
  - The PK of supertype is also PK for all subtypes
  - Each subtype has a FK to the supertype
- Preferred option by far!
  - Redundancy is eliminated:
  - Name and DOB are stored only once
  - Adjusts well to hierarchies/lattices

Person			Employee			Student		
SSN	Name	DOB	SSN	Department	Salary	SSN	GPA	StartDate
1234	Mary	1950	1234	Accounting	35000	1234	3.5	1997

# Specialization in SQL DDL

```
CREATE TABLE Person (  
    SSN INT PRIMARY KEY,  
    Name VARCHAR NOT NULL,  
    DOB DATE NOT NULL  
);
```

Person		
SSN	Name	DOB
1234	Mary	1950

Employee		
SSN	Department	Salary
1234	Accounting	35000


Student		
SSN	GPA	StartDate
1234	3.5	1997

```
CREATE TABLE Employee (  
    SSN INT PRIMARY KEY REFERENCES Person,  
    Department VARCHAR NOT NULL,  
    Salary INTEGER NOT NULL  
);
```

```
CREATE TABLE Student (  
    SSN INT PRIMARY KEY REFERENCES Person,  
    GPA REAL NOT NULL,  
    StartDate DATE NOT NULL  
);
```

# Categorization

- **Grouping of otherwise unrelated entities**
- New entity is a union = u in the circle
  - Can be total or partial = p/t on the line
- Can be checked with triggers similarly to specialization
- Notice the direction of the arc to the union
  - Refers to flow of “inheritance” – or which comes first...
  - Arcs matter → need to draw them (in the exam)!




# Categorization in SQL DDL

- New table for the categorical entity
  - New abstract PK with INTEGER / IDENTITY
- Add the PK as attribute to the other entities
  - With FK to the categorical entity
- Why is this important?
  - Sometimes AccountHolder participates in relationships...

```
CREATE TABLE AccountHolder (  
    AcctID INT PRIMARY KEY  
);
```

```
CREATE TABLE Person (  
    PNR INT PRIMARY KEY,  
    pname VARCHAR NOT NULL,  
    AcctID INT [NOT NULL]  
    REFERENCES AccountHolder (AcctID)  
);
```

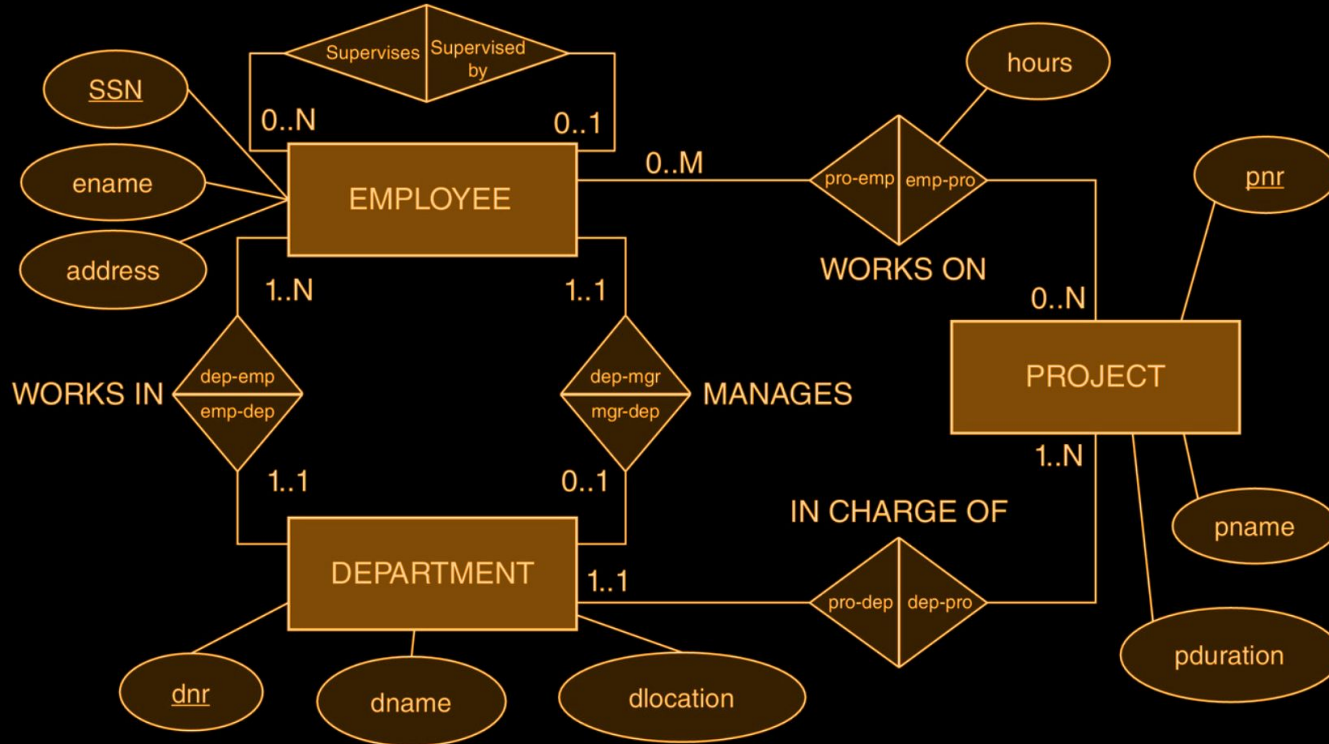
NOT NULL = Total  
NULL = Partial



# -- TODO -> DONE

- ✓ Conceptual Data Modeling (ER Diagram)
  - ✓ Entities and Attributes
  - ✓ Relationships
    - ✓ Cardinalities
    - ✓ Partial Relationship Keys
  - ✓ Weak Entities
  - ✓ Aggregation
  - ✓ Generalization/Specialization
  - ✓ Categorization
- ✓ Translation to SQL DDL

# Exercise: How many tables? What are the keys?



# Requirements to ERD

- Nouns = entities
  - Descriptive elements = attributes
- Verbs = relationships
  - Descriptive elements = attributes
  - Look for words implying participation constraints
  - No words  $\rightarrow$  0..N
- Example:
  - Professors have an SSN, a name, an age, a rank, and a research specialty.
  - Projects have a project number, a sponsor name (e.g., NSF), a starting date, ...
  - Each project is managed by one professor (1..1 on profs, 0..N on projs)
  - Professors may work on many projects (0..\* on both sides)
  - Each project must be reviewed by some professors (1..N on profs, 0..N on projs)

# Dealing with very large ER Diagrams

- Method 1: Very large paper!
  - One diagram with all the details
- Method 2: Outline + Details
  - One diagram with main entities and their relationships
  - One diagram per entity with attributes and weak entities
- Method 3: Components
  - Break the model into components
  - Details inside components
  - Some edge entities are repeated (details in one place)



# Limitations of ER Design

- ER diagrams do not capture all design details
  - Example: Multiple candidate keys
  - Must note missing details somewhere!
- Some aspects do not map well to SQL DDL
  - Example: 1..M cardinalities
  - Triggers (Lecture 4) can be used to handle some problems
  - Normalization (Lecture 6) provides a mechanism for fixing some problems
  - Some must simply be noted and addressed in code or ignored!

# Summary of Notation Extensions

- Relationship roles are generally unnecessary
  - No need to label roles
  - Except for unary relationship types
  - May put relationship name inside rhombus
  - Except for unary relationship types
- We allow partial keys of relationships
  - Underlined relationship attribute
  - Part of the PK of the resulting relationship relation
- Aggregation entity may cover only the relationship
  - Much easier to read!
- Allow 0..\* in place of 0..N/M/L
  - ... or simply use 0..N everywhere!

# Takeaways

one-to-zero  
one-to-one  
one-to-many  
many-to-many

## ER diagram captures entities and relationships

- Weak entities, specialization, categorization and aggregation allow capturing more detailed model characteristics
- This is hard – but also useful – so you must practice!

## Conversion to SQL DDL

- Entities and relationships mapped to relations
- Essentially an algorithmic process (with some options)
- This is hard – but also useful – so you must practice!

## Notation: MANY VARIANTS EXIST!

- We use the one from the book (as extended in lecture)
- You must use this notation in the homework and exam!!!

**Next Time in IDBS...**

# Introduction to Database Systems

## IDBS - Fall 2024

### Lecture 6

### Normalization

---

Readings: PDBM 6.2-6.4

Eleni Tziritza Zacharitou